YUANBO PANG

Agent Evaluation & Post-Training Measurement

yuanbopang@berkeley.edu · (+1) 408-460-7526 · yuanbopang.pro · GitHub · LinkedIn

Summary

Co-first author of **Agents' Last Exam** (ALE): I build the task environments, automated evaluators, and verifiable-reward pipelines used to measure frontier agents, and run those evaluations with frontier model teams — analyzing trajectories, investigating surprising results, and reporting failure modes back to the labs.

Education

University of California, Berkeley — B.S. EECS, GPA 3.762/4.0 (in-major 3.811) *Expected Dec 2026*

Stanford University (Summer Session) — CS 229: Machine Learning *Summer 2025*

Experience

Research Assistant, Berkeley RDI / BAIR (Prof. Dawn Song) *Oct 2025 – Present*

Agents' Last Exam (ALE) — co-first author arXiv:2606.05405 agents-last-exam.org

- Co-designed and built ALE's evaluation system: 1K+ expert-designed, long-horizon agent tasks across 55 subdomains, each packaged as an executable environment (`load`/`start`/`evaluate`) with programmatic graders — no hand labels in the scoring loop.
- Led the task pipeline and auto-scoring: cross-domain evaluators (hash, tabular, geometric, visual, behavioral) producing **verifiable rewards** usable for both measurement and RL post-training (RLVR); designed checks that keep graders trustworthy as tasks and models change.
- Operate large-scale sandboxed evaluation on GCP (container/VM task runners), including computer-use tasks with browser/desktop automation; run ALE evaluations with frontier model teams (OpenAI, Anthropic, Meta TBD Lab, Qwen, Seed; under lab guidelines) — analyzing agent trajectories, getting to the bottom of surprising results, and reporting failure modes back to the labs.
- For ALE V2, building ultra-long-horizon tasks from large-scale production software, with a harness that evaluates *teams* of agents: coordination, division of labor, and context handoff beyond single-agent benchmarks.

Robotics Research, Prof. Masayoshi Tomizuka's Lab, UC Berkeley *2026 – Present*

Robot-learning data benchmark — measurement methodology

- Score dataset quality by downstream finetune uplift — the Δ on each model's native benchmark (LIBERO, SimplerEnv, CALVIN) rather than proxy metrics; validated that rankings stay consistent across ACT, Diffusion Policy, OpenVLA, π_0, and GR00T.

Founder, Equile — The Dry Lab for AI Scientists *Apr 2026 – Present*

- Building complete, reproducible computing environments where research agents run real computational studies (trusted scientific codes, scheduled cluster runs), with a provenance record for every result — the same problems as operating an eval fleet: reproducible environments, job scheduling, live monitoring, trustworthy outputs.

Publications

Agents' Last Exam. Yiyou Sun*, Xinyang Han*, Weichen Zhang*, **Yuanbo Pang***, et al., Dawn Song. *Under review at NeurIPS 2026.* arXiv:2606.05405 #1 Paper of the Week, Hugging Face Daily Papers (Jun 2026). (*co-first author)

ScienceIDE: Turning World's Scientific Codebase into Agent Learnable Environments. Hejia Geng*, Zesen Huang*, ..., **Yuanbo Pang**, et al. arXiv:2609.19134 #2 Paper of the Day, Hugging Face Daily Papers (Sep 2026).

Agent's GameDev Exam. Zonglin Di, Yinhe Wang, **Yuanbo Pang**, Yiyou Sun, et al., Dawn Song. *Under double-blind review at ICLR 2027* — agent evaluation in game development environments; author version available upon request.

Skills

Evaluation: benchmark & task design, automated evaluators/verifiers, verifiable rewards (RLVR), agent trajectory analysis, sandboxed task environments, regression investigation

Engineering: Python (primary), PyTorch, TypeScript, Docker, GCP, Modal, Postgres, agent harnesses, tool-use & MCP, Linux